Week 1: Data Collection for Machine Learning

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra

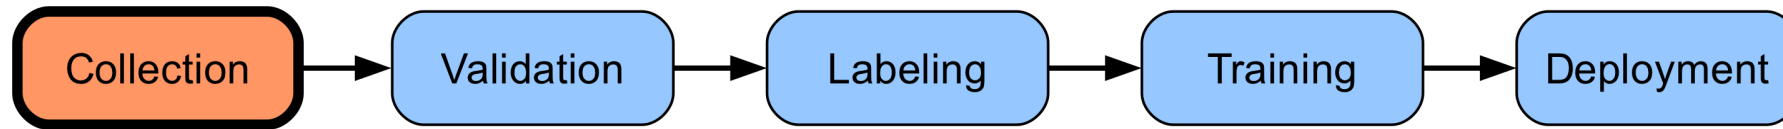IIT Gandhinagar

The Netflix Movie Recommendation Problem

Scenario: You work at Netflix as a data scientist.

The Task: "Predict which movies will be successful to decide our next acquisitions."

The Bottleneck: We have no data.

Today's Focus: How do we build the dataset to solve this problem?

The ML Pipeline

Collection → Validation → Labeling → Training → Deployment

diagram-generators/data_pipeline_flow.py

Garbage In, Garbage Out:

- 80% of ML work is data engineering.

- Sophisticated models cannot fix broken data.

- **Goal**: Automate the collection of high-quality data.

Data Sources Strategy

We need features: *Title, Budget, Revenue, Reviews, Cast*.

Source Type	Example	Pros	Cons
Public APIs	OMDb, TMDb	Structured, Reliable	Rate limits, Cost
Web Scraping	IMDb, Rotten Tomatoes	Free, Flexible	Fragile, IP bans
Datasets	Kaggle, Hugging Face	Clean, Ready	Static, Generic

Plan: Use OMDb API for base data + Scraping for reviews.

Part 1: The Web Protocol (HTTP)

How browsers (and scripts) talk to servers.

Client-Server Architecture
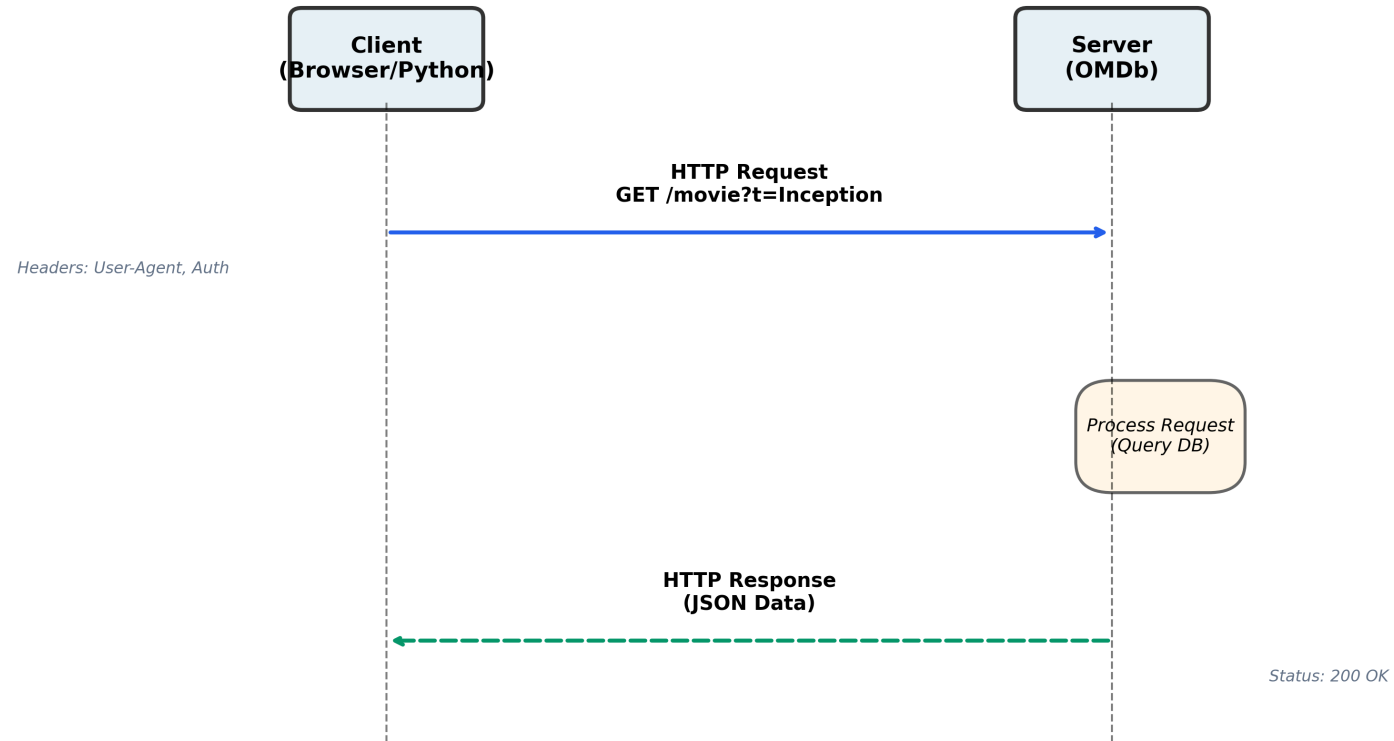
diagram-generators/http_request_sequence.py

6

Anatomy of a Request

URL: `https://api.omdbapi.com/?apikey=123&t=Inception`

Methods:

- `GET` : Retrieve data (Read). *Safe, Idempotent.*
- `POST` : Send data (Write). *Creating resources.*
- `PUT` : Update data.
- `DELETE` : Remove data.

Status Codes:

- `200` : OK.
- `400` : Bad Request (Client error).
- `401/403` : Unauthorized (Check your API Key).

Part 2: CLI Tools (curl & jq)

Test APIs before writing code.

curl: The HTTP Swiss Army Knife

Fetch data:

```
curl "http://www.omdbapi.com/?apikey=$KEY&t=Inception"
```

Inspect headers (`-I`):

```
curl -I "https://google.com"
# HTTP/2 200
# content-type: text/html
```

Why use curl?

- Language agnostic.

- Instant debugging.

- "Copy as curl" from Chrome DevTools.

9

jq: JSON Processor

Raw JSON is unreadable. `jq` makes it useful.

Pretty print:

```
curl ... | jq
```

Filter fields:

```
# Get just the title and rating
curl ... | jq '{Title, imdbRating}'
```

Filter array elements:

```
# Get titles of movies created after 2010
cat movies.json | jq '.[] | select(.Year > 2010) | .Title'
```

Part 3: Python `requests`

Automating the process.

The Synchronous Pattern

`requests` is **blocking**. The program stops until the server responds.

```python
import requests

def get_movie(title):
    url = "http://www.omdbapi.com/"
    params = {"apikey": "SECRET", "t": title}

    try:
        # Block here until response arrives
        resp = requests.get(url, params=params)
        resp.raise_for_status() # Check for 4xx/5xx
        return resp.json()
    except Exception as e:
        print(f"Failed: {e}")
        return None
```
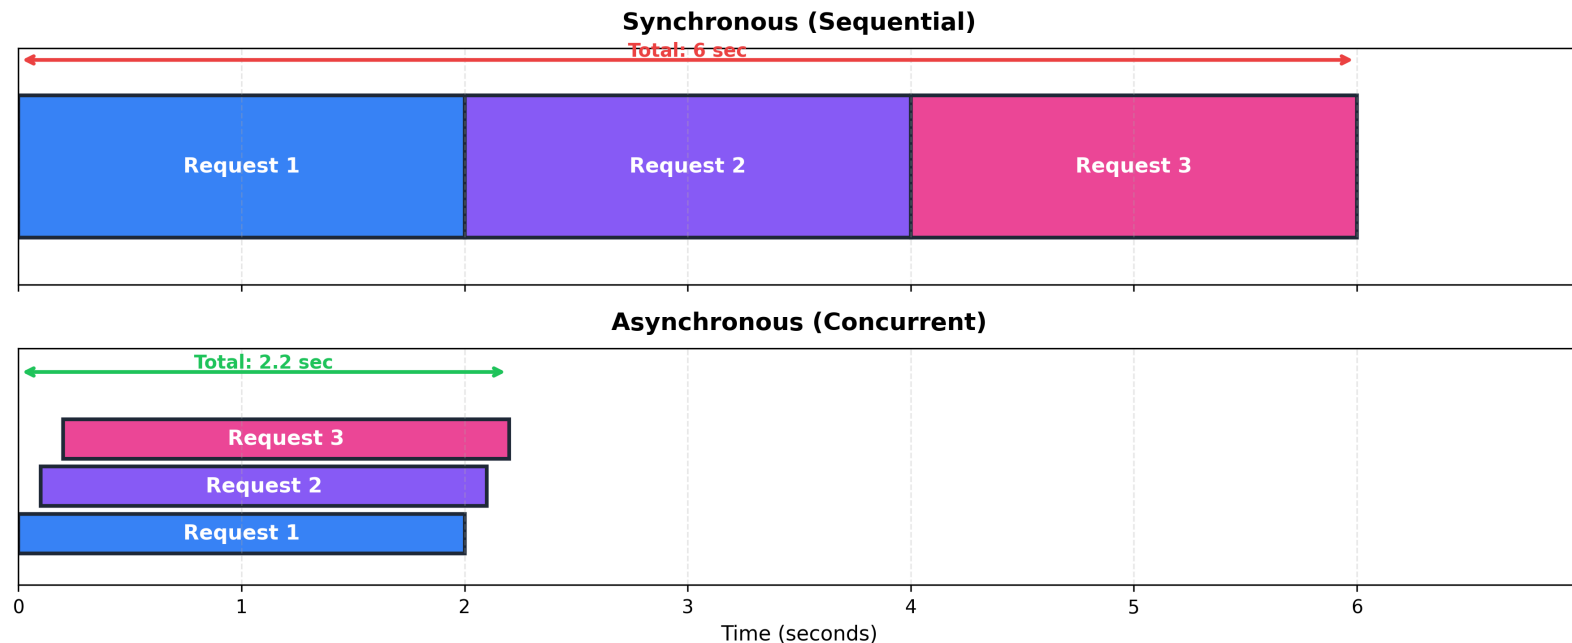
12

Advanced: Async IO (Conceptual)

Problem: Fetching 1,000 movies sequentially is slow.

Solution: Asynchronous Requests (`aiohttp`, `httpx`).

diagram-generators/sync_vs_async_timing.py

We will implement Async in Week 10 (FastAPI)

Handling Rate Limits

Exponential Backoff: If failed, wait longer.

```python
import time

def fetch_with_retry(url, retries=3):
    for i in range(retries):
        resp = requests.get(url)
        if resp.status_code == 429: # Rate limit
            wait = 2 ** i  # 1s, 2s, 4s...
            time.sleep(wait)
            continue
    return resp
```

Ethical Scraping:

- Always check `robots.txt` (e.g., `imdb.com/robots.txt`).
- Identify your bot (`User-Agent: MyStudentBot/1.0`).

Part 4: Web Scraping (BeautifulSoup)

When there is no API.

The DOM Tree

HTML is a tree of tags. `BeautifulSoup` traverses it.

```
<div class="movie-card">
    <h1>Inception</h1>
    <span class="rating">8.8</span>
</div>
```

Parsing Logic:

1. Find container: `div.movie-card`

2. Extract children: `h1` (text), `span.rating` (text).

Scraping Strategies

Strategy	Library	Use Case
Static HTML	`BeautifulSoup`	Simple sites, fast, lightweight.
Dynamic JS	`Playwright`	Single Page Apps (React/Vue), infinite scroll.
Visual	`GPT-4V`	Complex layouts, CAPTCHAs (expensive).

Note: Playwright is covered in the Lab.

Data Licensing & Ethics

Can I use this data?

1. **Public Domain (CC0)**: Free to use.

2. **Attribution (CC-BY)**: Use, but credit source.

3. **Non-Commercial (NC)**: Academic use OK, selling NOT OK.

4. **Copyright**: "All Rights Reserved".
   - *Fair Use*: Small excerpts for research *might* be okay.
   - *Scraping*: Generally legal for public data (US hiQ v LinkedIn), but ToS violation can get you banned.

Rule: Be respectful. Don't DDoS the server.

Summary

1. **APIs > Scraping**: Always look for an API first (stable, legal).

2. **Tools**: `curl` for quick checks, `requests` for scripts.

3. **Robustness**: Handle errors, retries, and rate limits.

4. **Ethics**: Respect `robots.txt` and server load.

Next Up: Now that we have data, it's probably messy. **Week 2: Data Validation**.