

Deployment on Constrained Devices

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra, IIT Gandhinagar

The "Edge" Challenge

Scenario: You trained a ResNet-50. It's 100MB.

You want to run it on a Raspberry Pi or a Mobile Phone.

Constraints:

- 1. Memory:** Device has 2GB RAM, model needs 4GB.
- 2. Latency:** Inference takes 5s, user needs <100ms.
- 3. Power:** GPU drains battery in 20 mins.
- 4. Storage:** App limit is 50MB.

Solution: Model Optimization.

Techniques Overview

```
graph TD A[Trained Model] --> B[Quantization]; A --> C[Pruning]; A --> D[Knowledge Distillation]; A --> E[Architecture Search]; B --> F[Optimized Model]; C --> F; D --> F; E --> F;
```

Today's Focus:

- 1. Quantization:** Lower precision math.
- 2. Pruning:** Removing useless connections.
- 3. ONNX:** Efficient Runtime.

Quantization: Theory

Standard Training: Float32 (32-bit floating point).

Quantization: Convert to Int8 (8-bit integer).

Formula:

$$Q(x) = \text{round} \left(\frac{x}{S} + Z \right)$$

- S : Scale
- Z : Zero-point

Impact:

- **Size:** 32 bits -> 8 bits = **4x reduction**.
- **Speed:** Integer math is faster than float math on CPUs.
- **Accuracy:** Minimal drop (<1%) for robust models.

Types of Quantization

1. Post-Training Quantization (PTQ):

- Train normal Float32 model.
- Calibrate with small dataset.
- Convert to Int8.
- *Easiest.*

2. Quantization-Aware Training (QAT):

- Simulate quantization *during* training.
- Model learns to adapt to lower precision.
- *Best Accuracy.*

Pruning: Theory

Idea: Neural Networks are over-parameterized. Many weights are near zero.

Action: Set small weights to exactly zero.

Structured vs Unstructured:

- **Unstructured:** Random zeros. Good for compression, bad for speed (sparse matrices need hardware support).
- **Structured:** Remove entire channels/filters. Good for speed (smaller matrix).

ONNX: Open Neural Network Exchange

The Universal Bridge

```
graph LR A[PyTorch] --> D[ONNX Graph]; B[TensorFlow] --> D; C[Scikit-Learn] --> D; D --> E[ONNX Runtime (ORT)]; E --> F[Android]; E --> G[Raspberry Pi]; E --> H[Browser (WASM)];
```

Why use it?

- **Interoperability:** Train in PyTorch, deploy in C++.
- **Optimization:** ORT applies graph fusions (e.g., Conv+ReLU merging).

Lab Preview

Hands-on Optimization:

1. **Baseline:** Measure size/speed of ResNet-18.
2. **Pruning:** Use `torch.nn.utils.prune` to remove 30% of weights.
3. **Quantization:** Apply PyTorch dynamic quantization.
4. **ONNX Export:** Convert and run with ONNX Runtime.

Goal: Make the model 2x faster and 4x smaller!